

Routing in Delay-Tolerant Networks Comprising Heterogeneous Node Populations

Thrasyvoulos Spyropoulos, *Member, IEEE*, Thierry Turletti, *Member, IEEE*, and
Katia Obraczka, *Member, IEEE*

Abstract—Communication networks are traditionally assumed to be connected. However, emerging wireless applications such as vehicular networks, pocket-switched networks, etc., coupled with volatile links, node mobility, and power outages, will require the network to operate despite frequent disconnections. To this end, *opportunistic* routing techniques have been proposed, where a node may store-and-carry a message for some time, until a new forwarding opportunity arises. Although a number of such algorithms exist, most focus on relatively homogeneous settings of nodes. However, in many envisioned applications, participating nodes might include handhelds, vehicles, sensors, etc. These various “classes” have diverse characteristics and mobility patterns, and will contribute quite differently to the routing process. In this paper, we address the problem of routing in intermittently connected wireless networks comprising multiple classes of nodes. We show that proposed solutions, which perform well in homogeneous scenarios, are not as competent in this setting. To this end, we propose a class of routing schemes that can identify the nodes of “highest utility” for routing, improving the delay and delivery ratio by four to five times. Additionally, we propose an analytical framework based on fluid models that can be used to analyze the performance of various opportunistic routing strategies, in heterogeneous settings.

Index Terms—Delay-tolerant network, intermittent connectivity, routing, utility, replication, fluid model.

1 INTRODUCTION

THE traditional view of networks as a connected graph over which end-to-end paths need to be established might not be appropriate for modeling existing and emerging wireless networks. Due to wireless propagation phenomena, node mobility, low-power nodes periodically shutting down and waking up, etc., connectivity in many wireless networks is, more often than not, intermittent. Despite this limited and/or episodic connectivity, many emerging wireless applications could still be supported. Some examples are the low-cost Internet provision in remote or developing communities [1], [2], [3], vehicular networks (VANETs) for dissemination of location-dependent information (e.g., local ads, traffic reports, parking information, etc.) [4], “pocket-switched” wireless networks to extend and sometimes bypass access point connectivity to the Internet [5], [6], [7], tactical networks operating in an intermittent fashion for low probability of interception and low probability of detection (LPI/LPD) reasons [8], underwater networks [9], etc.

These networks are often referred to collectively as Delay-Tolerant Networks (DTNs) [10] and target scenarios where: 1) providing guaranteed connectivity is very challenging due to the particular environment in hand

(e.g., deep space, underwater) and 2) providing *some* level of service, even if more limited than what we are accustomed to in the wired Internet, may present significant economical, social, and/or cultural advantages to the end user or operator [1], [6], [11]. The goal is to provide certain networking services, despite frequent and long-lasting disconnections. As such, the targeted services are usually ones that are asynchronous and can thus often tolerate some amount of delay with little perceived degradation of quality (e.g., e-mail and cached/offline Web access [12], sensor data gathering [13], service or information discovery [6], etc.).

However, networking under such intermittent connectivity is particularly challenging, as many of the assumptions made by traditional protocols (TCP, DNS, etc.) “break” in this context [10]. Arguably though, routing is one of the biggest hurdles to overcome. Traditional routing protocols, both table-driven or proactive (e.g., link-state-based routing protocols like OSPF and OLSR) and reactive ones (e.g., DSR, AODV), assume the existence of a complete end-to-end path, and try to discover it *before* any useful data are sent. As a result, their performance deteriorates drastically as connectivity becomes increasingly sporadic and short-lived.

To overcome this problem, researchers have proposed a variety of *opportunistic* routing schemes [14], [15], [16] most of which exhibit the following basic characteristics: 1) a message may be *stored* and *carried* by a node for long periods of time until a communication opportunity arises (“mobility-assisted”); 2) local forwarding decisions are made independently, and often with little or no knowledge about the destination’s position, with the goal that a message will *eventually* be delivered (“opportunistic”); and 3) multiple copies of the same message may be propagated in parallel (“replication”), to increase the probability of at least one being delivered.

• T. Spyropoulos is with ETH Zürich, Inst.f.Techn.Informatik u.Kommunik.Netze, ETZ G 96, Gloriastrasse 35, 8092 Zurich, Switzerland. E-mail: spyropoulos@tik.ee.ethz.ch.

• T. Turletti is with INRIA, 2004 route des Lucioles, B.P. 93, 06902 Sophia Antipolis, France. E-mail: turletti@sophia.inria.fr.

• K. Obraczka is with the Computer Engineering Department, Jack Baskin School of Engineering, University of California, Santa Cruz, 1156 High Street, Santa Cruz, CA 95064. E-mail: katia@cse.ucsc.edu.

Manuscript received 23 Oct. 2007; revised 23 July 2008; accepted 11 Sept. 2008; published online 18 Dec. 2008.

For information on obtaining reprints of this article, please send e-mail to: tmc@computer.org, and reference IEEECS Log Number TMC-2007-10-0318. Digital Object Identifier no. 10.1109/TMC.2008.172.

One of the schemes proposed that has shown promising performance is that of “Controlled Replication” or “Spray and Wait” [16], [17], [18]. This scheme distributes a *small* and *controlled* number of message replicas to the *first few nodes encountered*, then each of the nodes that received a copy in this first phase is only allowed to give it to the destination itself. This approach exhibits a number of desirable characteristics: 1) it consumes much fewer resources than epidemic routing [19] and its variations [20], [13], where potentially all nodes receive the message; 2) due to its simplicity, it can be *controlled* to achieve the desired performance trade-off (resources used versus delivery delay/probability), even in an almost unknown environment, by appropriately choosing the number of copies [16]; and 3) it has been recently shown to be quite robust to some security attacks [21]. The rationale behind the protocol is that picking just a few relays, even randomly, could create enough redundancy to ensure that one will encounter the destination soon. If most nodes in the network are highly “mobile,” roughly homogeneous in terms of their message delivery capability and make decisions independently, there is indeed little or no penalty on the delivery delay or delivery probability [18], [16], despite the greedy choice of relays and their small number. In fact, it has been proven that if nodes are homogeneous and mobility is stochastic and IID, handing over copies *greedily* to the first nodes encountered is optimal [17], [22].

Yet, in many scenarios, the requirement for nodes with homogeneous capabilities and independent mobility does not hold. For example, imagine a scenario where the majority of nodes tend to spend most of their time around the same nodes (e.g., employees in the same office or floor, or monitored animals in the same herd [13]), while only a small number of nodes tend to move often between disconnected parts of the network (e.g., vehicles, nodes with a more “social” daily routine, etc.). In such a scenario, handing over copies to the *first few nodes encountered* implies that some or all of these copies will end up with nodes that may never encounter the destination. In general, *when nodes in a network are heterogeneous in terms of their ability to deliver a message, greedy replication may fail to discover the “better” relays, especially if the latter are much less common than the “nonuseful” ones.*

To address this problem, some existing schemes allow an initially chosen relay to hand over its message replica to a “better” relay. This is usually achieved by maintaining some utility or fitness function by nodes, which is then used to perform a gradient-based search for the destination [15], [16], [23]. However, this implies that each message replica may be transmitted more than once or even loop around the network. Although proposed solutions have not been shown to exhibit a significant increase in overhead, *on average* (see, e.g., [16]), the *worst-case* performance can be arbitrarily bad in principle. What is more, unlike the case of controlled replication, it is significantly more difficult to analytically predict the number of transmissions and expected delay for this scheme [16]. This can be a rather undesirable feature of the protocol for some applications. For example, assume that an application has stringent energy constraints (bounded amount of energy per message) but some flexibility in the delivery ratio or delay (e.g., sensor data with inherent redundancy). As a different example, in a self-organized network where a credit-based

system is used to incite collaboration between nodes [24], it is important for the user to be aware of exactly how many credits a message will consume. Summarizing, although allowing each copy to be forwarded over multiple hops can sometimes reduce the message delivery delay, it may also result in unpredictable and uncontrolled resource usage.

For these reasons, in this paper, we focus our study on strategies which use a fixed number of copies and no forwarding, similar to the basic greedy replication algorithm. Our aim is to maintain control and predictability of resource usage as our primary objective, while at the same time, try to achieve the best performance possible in terms of delay and delivery ratio. At the same time, we explicitly assume that our network consists of highly heterogeneous nodes and require a routing protocol that can efficiently address this. In short, we try to answer the following question: *given a fixed budget of message copies, how can we best allocate it to a network consisting of nodes with varying capabilities and behaviors?*

To this end, we propose the idea of *smart replication* or *utility-based replication*, where instead of naively (greedily) handing copies to the first nodes encountered, we choose relays according to some utility function. We examine in detail situations where nodes might be heterogeneous in one or both of the following aspects: 1) in terms of which other nodes they tend to encounter often (e.g., due to social links, colocation, etc.) and 2) in terms of generic characteristics according to which nodes can be categorized into various “node classes” (e.g., static nodes, vehicles, pedestrians, etc.) with relatively homogeneous characteristics (e.g., mobility, resources, etc.) inside the class, but largely heterogeneous between classes. We propose and evaluate a number of simple, yet efficient, replication strategies, and also examine how to “tune” important parameters of each protocol in practice. Using both synthetic and trace-driven simulations, we show that utility-based replication schemes consistently outperform greedy replication, by up to four to five times in the scenarios considered. This work extends on our work described in [25], where we have performed a preliminary presentation and simulation-based evaluation of some replication strategies using synthetic simulations only.

Additionally, we propose an analytical framework based on fluid models to analyze the performance of mobility-assisted routing schemes in heterogeneous networks consisting of different populations of nodes. Fluid models are quite popular in the mathematical biology community to model epidemics in populations [26]. Although an initial effort to apply such models in a DTN context has been made in [27], it was done so for homogeneous networks only. To the best of our knowledge, this is the first work that analyzes encounter-based routing protocols in networks consisting of heterogeneous node classes. We use this framework to analyze the performance of greedy- and utility-based replication, as well as the performance of epidemic routing [19], a popular routing protocol for disconnected environments, which can be optimal under certain, ideal conditions. We believe that our model could provide a better understanding of various encounter-based processes (e.g., electronic virus spread, DTN podcasting [28]) in real, heterogeneous wireless networks.

In the next section, we go over related work. Then, in Section 3, we define the general utility-based replication

algorithm and describe three specific strategies based on it. In Section 4, we present synthetic and trace-driven simulation results comparing these three utility-based schemes against greedy replication, and then, in Section 5, we give some insight into how utility-based replication can be tuned in practice. In Section 6, we develop our theoretical model and use it to derive the delay of various mobility-assisted protocols. Finally, we conclude and present some future work directions in Section 7.

2 RELATED WORK AND CONTRIBUTIONS

A number of (mobility-assisted) routing protocols have been proposed for DTN networks that assume known or enforced future connectivity (e.g., [2], [29]). If future connectivity is intermittent but known, existing routing algorithms (e.g., Dijkstra's) can be appropriately modified to discover shortest paths-over-time. In general, problems with such future knowledge can often be formulated as dynamic multicommodity flow or quickest transshipment problems [30], [31]. A couple of interesting extensions of this work is to try to discover and maintain the shortest paths-over-time, under uncertain or probabilistic information about future connectivity [32] or to employ hybrid MANET-DTN routing solution, when there is some stability or connectivity in certain parts of the network [33], [34].

Similarly to [15], [16], [19], [27], we, in this paper, follow a different approach assuming random connectivity (note: random does not necessarily mean unpredictable) and "opportunistic" forwarding decisions. A number of such opportunistic routing proposals are described in [14]. Among these protocols, controlled replication variants have been found to present a desirable trade-off in many scenarios, especially when it comes to efficient resource usage [27]. However, it has also been shown that simply generating and handing over a few redundant copies may not often suffice in situations where the mobility or interactions between nodes are highly correlated and follow specific patterns [16], [35].

When such "structure" exists in the mobility patterns, it has been demonstrated that trying to infer these patterns and use the knowledge to probabilistically predict future connectivity can lead to better forwarding decisions and improve performance [16], [20], [23], [36], [37]. However, this is often done assuming a specific utility function and a specific mobility setting, while, at the same time, not providing explicit control over the resources used per message. In this paper, we aim to explicitly combine resource control and use of history-based information into an "online" allocation problem. Interesting efforts to formulate DTN routing as a resource allocation problem can also be found in [38], but with reliability as the goal, and in [39], but without explicit control on the number of message replicas used and a focus mostly on homogeneous settings (e.g., bus network).

Furthermore, the delay of various encounter-based or mobility-assisted routing schemes has been analyzed recently, for example, epidemic routing under no contention [15], [27], [40], [41] and under contention [42], single-copy schemes [15], greedy replication [17], [22], and other flooding-based schemes [27]. These efforts either use stochastic models (e.g., transient analysis of random walks

on graphs [15], [16] or Markov Chains [40], [41]) or deterministic approximations of these (*fluid models*) [27]. However, the common assumption in all these works is that all nodes move in an IID manner. Although this often simplifies analysis and allows, for example, simple Markov Chain models to be used, nodes in real life (and their respective mobility behaviors) are more often heterogeneous. For this reason, we decide to relax the assumption of identical mobility and try to improve existing analytical models for epidemic message spreading, by allowing nodes to belong to different *mobility classes*.

As a final note, a large body of work on modeling of epidemic processes exists in the areas of Epidemiology and Mathematical Biology. Recently, it has been recognized that popular models from these communities can have useful application to encounter-based routing protocols also [17], [27]. Furthermore, humans are generally heterogeneous in terms of their susceptibility to different diseases (e.g., male vs. female, children vs. elderly, etc.) and form societal structures (e.g., families, communities, schools, etc.) inside which spreading characteristics might be largely different than in the general population. As a result, effort has also been devoted to model spreading of epidemic diseases in heterogeneous populations [43], stratified populations [44], [45], and complex social networks (e.g., small-world networks [46]). However, the focus of these works is on whether an epidemic will spread or not (e.g., "epidemic thresholds") and on limiting behaviors, in general (e.g., convergence points). In the case of encounter-based message forwarding, we are interested more in the transient behavior of the process and particular *stopping times* [47], like, for example, "the time until node X gets infected." To our best knowledge, such quantities have not been widely studied in the context of heterogeneous networks in either the Epidemiology or the Networking community.

3 UTILITY-BASED REPLICATION

3.1 Greedy Replication in Heterogeneous Environments

Controlled replication or "Spraying" is considered to be an efficient method to reduce the large overhead of epidemic-based schemes without often incurring significant delay penalties [16], [17], [18]. The basic controlled replication algorithm is as follows.

Definition 3.1 (Spray and Wait or Controlled Replication).

When a new message is generated at a source node, this node also creates L "forwarding tokens" for this message. A forwarding token implies that the node that owns it can spawn and forward an additional copy of the given message. During the spraying phase, messages get forwarded according to the following rules:

- *If a node (either the source or a relay), carrying a message copy and $c > 1$ forwarding tokens for this message, encounters a node with no copy of the message,¹ it spawns and forwards a copy of that message to the second node; it also hands over $l(c)$*

1. We assume that a message ID vector exchange similar to that performed in epidemic routing occurs [19].

tokens to that node ($l(c) \in [1, c - 1]$) and keeps the rest $c - l(c)$ for itself (“Spray” phase).

- When a node has a message copy and $c = 1$ forwarding tokens for this message, then it can only forward this message to the destination itself (“Wait” phase).

There are two flavors of the basic scheme that have been proposed. In the 2-hop version of the scheme [41], only the source may forward an extra copy (i.e., $l(c) = 1$ in the above description). In the (binary) tree-based version [16], [17], $l(c) = \lfloor \frac{c}{2} \rfloor$ for any node with $c > 1$ tokens. The common characteristic between these two algorithms is that they are both greedy. By “greedy” here we mean that any opportunity to forward one of the available copies to a new node with no copies will always be seized upon. In other words, the budget of L message replicas will be distributed to the L first nodes encountered.

The tree-based greedy algorithm is shown to be optimal in a homogeneous environment (IID node mobility) [16], [17]. The intuition behind this is that if all nodes are statistically equivalent [48], there is no benefit in waiting when a forwarding opportunity arises. Nevertheless, if nodes are heterogeneous (e.g., different mobility characteristics, different resources, etc.), greedy distribution of message copies can be easily shown to be suboptimal. Consider, for example, an intermittently connected network of mobile nodes, where a percentage p of the nodes is not (that) useful in delivering a message to a remote destination (e.g., move very slowly, or encounter few new nodes over time). We will call these “snail” nodes. If we assume uniform mixing rates [48] for both normal and snail nodes, then with greedy replication pL of the copies, on average, would end up with snail nodes.² Using the delay equations of [16] (e.g., (2)) or [17], it is not difficult to show that the performance degradation compared to the case where all nodes are “nonsnail” ones is approximately given by

$$\frac{\text{Delay}(p \text{ snail})}{\text{Delay}(\text{normal})} \approx c_1 + \frac{c_2}{1 - p},$$

where c_1, c_2 are constants such that $c_1 + c_2 = 1$. (The first term of the sum is related to the spraying phase, while the second term to the wait phase, with c_1, c_2 absorbing the various quantities common to both cases, e.g., number of copies, spray phase duration, etc.)

This equation implies that if the number of snail nodes is large, then the delay of greedy Spray and Wait quickly increases. The reason for this is that greedy replication cannot identify nodes that are not useful, and mistakenly hands them over some copies. The larger the percentage of nonuseful relays in the network, the larger the negative impact of greediness. Of course, a more appropriate comparison would be against the delay of the optimal algorithm given p snail nodes. We defer this to Section 6. Nevertheless, the above equation already provides enough insight on the expected performance degradation due to heterogeneity.

3.2 Utility-Based Spraying

Based on the previous exposition, we can draw the following conclusion: in a heterogeneous environment,

2. This is not exactly true: If snail nodes move less frequently around the network than normal ones, it may be the case that they are also encountered by other nodes less frequently; in that case, pL would just be an upper bound.

where a limited budget of L message copies needs to be distributed to L relays, a mechanism is necessary to distinguish the “better” relays and avoid using the least useful ones. Ideally, we would like to find the L best relays in the network (given some optimization criterion). However, this problem is not trivial even in moderately complex scenarios [49], especially, given the fact that candidate relays appear (i.e., are encountered) not all together, but in an online fashion. Therefore, here, we will turn our attention to heuristic methods to decide on the fitness or utility of a given node as a relay.

Definition 3.2 (Utility-Based Spraying). Similar to the basic spraying algorithm (see Definition 3.1), Utility-based Spraying uses forwarding tokens to grant a node the right to further forward message copies. Additionally,

- Each node i maintains a utility function $U_i(j)$ for every other node in the network j . $U_i(j)$ reflects the probability that node i will deliver a message to node j , and it may be based on a number of different parameters (e.g., encounter history, mobility, friendship index with j , etc.).
- If a node i (either the source or a relay) carrying a message copy for a destination d and $c > 1$ forwarding tokens for this message encounters a node j with no copy of the message, it spawns and forwards a copy of that message to the second node according to one of the following rules:
 - **Rule 1.** If $U_j(d) > U_{th}$ for some U_{th} threshold value (absolute utility criterion).
 - **Rule 2.** If $U_j(d) > U_i(d)$ (relative utility criterion).
 It also hands $l(c) = \lfloor \frac{c}{2} \rfloor$ forwarding tokens to that node and keeps the rest $\lceil \frac{c}{2} \rceil$ for itself (i.e., tree-based).³

Each algorithm could use either of the forwarding rules above (“absolute utility” or “relative utility”) or a combination thereof (e.g., “use rule 1 to ensure a minimum utility, and then rule 2 among the nodes that qualify”). Rule 1 requires an appropriate threshold parameter U_{th} to be set in every case. Rule 2 is easier to implement, yet it does not guarantee that all “bad nodes” will be avoided (e.g., “very low utility” nodes could still give copies to “low utility” nodes).

We are going to describe different variations of the basic algorithm in terms of the utility function they use. Before we choose our utility function, we note that, in a real setting, there are a number of different reasons why a given node might be a “better” relay than another. These might range from special “bonds” with the destination to higher resources available to a node to higher reliability or trustworthiness of a node. In general, though, candidate utility functions could be broadly categorized into destination-dependent (“DD”) and destination-independent (“DI”) functions.

Destination-Dependent (DD) Utility. One node may be the best relay for one destination (d_1) and another node the

3. A more generic approach would be to make $l(c)$, the number of tokens forwarded, a function of the relay’s utility: $l_c = f(c, U_j(d))$. This would increase the flexibility but also the complexity of the distribution algorithm. What is more, in situations where not all nodes are honest, avoiding “overinvesting” on nodes advertising very high utilities might be preferable.

best relay for a different destination (d_2). In other words, for DD functions

$$U_i(d_1) > U_j(d_1) \text{ but } U_i(d_2) < U_j(d_2), \quad d_1 \neq d_2. \quad (1)$$

Examples of DD utility functions could be those based on age-of-last-encounter for a given destination [15], [23], social relation with a given destination [50], correlated mobility patterns [51], etc. Destination-dependent utility functions impose a larger overhead on nodes, as in essence, they need to maintain an entry for every other node in the network. To reduce this, overhead caching techniques could be used (e.g., storing only the m highest utility values).

Destination-Independent (DI) Utility. The “utility” of a given node is independent of any destination; instead, it depends on some special characteristic(s) this node has. This implies that one node may be the best relay for most or all destinations. In other words, for DI functions, it holds, in general, that

$$U_i(d_1) \geq U_j(d_1) \Rightarrow U_i(d) \geq U_j(d), \text{ for most or all } j, d. \quad (2)$$

Examples of nodes which are highly preferable as relays for any destination would be nodes with high and frequent mobility (e.g., vehicles), nodes with many “friends” (e.g., hubs in scale-free networks), or nodes with higher resources (e.g., a “message ferry” [14]). DI utility functions have a smaller overhead than DD functions as they require each node to maintain only a single utility value each. Yet, DI functions also imply that the “better” nodes might have to bear a higher forwarding overhead than others. This might result in poorer load-balancing and utilization of the total network capacity, and/or faster battery drainage of a few nodes. We address this issue in Section 5.

We will now turn our attention to specific replication strategies, some of which use DD utility functions and others DI ones. It is important to note that it is also possible to define hybrid algorithms that take into account both the general *fitness* of a node as well as destination-specific information when making a forwarding decision. However, it is beyond the scope of this paper to examine in depth all possible options. Instead, we are interested in describing a few simple replication heuristics, depending on the expected “structure” of the network and the types of nodes one would expect to encounter. We will demonstrate later, in Section 4, that even these simple solutions can have a significant impact on performance. Furthermore, it is our belief that finding an “optimal” utility function requires first to choose a particular target application. It is then that one should aim to optimize one of these simple strategies (or a combination of them) based on the specific application characteristics and requirements.

All algorithms, hereafter, are variations of the basic utility-based replication scheme described in Definition 3.2.

3.3 Last-Seen-First (LSF) Spraying

The basic idea is to choose as relays the nodes that have seen the destination most recently. Thus, LSF is an example of a destination-dependent utility function.

Definition 3.3 (LSF Spraying). *Each node i maintains a timer $\tau_i(j)$ for every other node j in the network, which records the time elapsed since the two nodes last encountered each other as follows: initially, set $\tau_i(i) = 0$ and $\tau_i(j) = \infty, \forall i, j$; if i*

encounters j , set $\tau_i(j) = \tau_j(i) = 0$; otherwise increase each $\tau_i(j)$ at every time unit. Finally, $U_i(j) = \frac{1}{1+\tau_i(j)}$.

The basic idea of using age-of-last-encounter timers has been introduced to advance a single message copy toward its destination in a connected mobile network, using gradient-based routing [23]. Unlike [23], [16], where one starts with a random set of relays and use age-of-last-encounter to find better and better ones, LSF uses it to choose a promising set of relay nodes from the beginning. Strategies based on last encounter could be efficient, for example, when nodes are separated into groups (e.g., work groups, departments, herds, etc.) and members in a group see each other more often than members out of it. Another example could be when certain nodes tend to see much larger subsets of nodes, in general, and thus would also encounter a randomly chosen destination more often than the average node (note that frequency of encounter can also be captured indirectly in the age-of-last-encounter from a statistical point of view).

3.4 Most-Mobile-First (MMF) Spraying

This is our first algorithm that uses a DI utility criterion. It is a simple scheme that assumes that some nodes are more “mobile” (or, in general, more “capable”) than others. This scheme would work well, for example, when the majority of nodes participating in the ad hoc network are pedestrians moving within a local area, while a few nodes are vehicles (e.g., buses, taxis, etc.) that tend to visit larger areas and with higher speeds (and thus, have a higher chance of meeting new nodes). For now, we will assume that each of these nodes carries a *label* that states the type of the node, e.g., “BUS,” “TAXI,” “PEDESTRIAN,” and “BASE STATION.” Although, in some scenarios, it would not be too burdensome to manually configure a label (e.g., by setting some software parameter when installing a radio, say, on a bus); in the next section, we discuss one way of assigning labels automatically.

Definition 3.4 (MMF Spraying). *Assume that there are m total node labels: $LABEL_1, LABEL_2, \dots, LABEL_m$. Each node i is assigned one of the labels, let $LABEL(i)$, and we define the utility of node i as $U_i(j) = U_i = LABEL(i), \forall j$. Finally, labels are put in a preference order (\succ):*

$$LABEL_1 \succ LABEL_2 \succ \dots \succ LABEL_m.$$

This assignment of preference order to labels can be made offline, based on the general mobility statistics and perceived usefulness of different types of nodes. Furthermore, one could also have different preference orders depending on the type (label) of the destination. For example, nodes of $LABEL_j$ may not be good relays in general, but excellent candidates if the destination is also of $LABEL_j$. A label assignment along these lines is performed in [52], where labels are based on affiliation. Thus, a node is a good relay for any destination with the same affiliation. This would actually bring the scheme somewhere in between DI and DD.

3.5 Most-Social-First (MSF) Spraying

In the previous scheme, we assume that, somehow, information about the mobility characteristics of a node is available. However, in many scenarios, the same wireless

device might be carried by a pedestrian, left at the office desk, or lie inside a vehicle at times. Hence, we need a mechanism that could estimate the “degree of mobility” *online*. What is more, some nodes might encounter more nodes than average not due to mobility, but just because they visit some “hub” locations (e.g., cafeteria) more often, or have more *social links* than average. This implies that a more appropriate metric of the utility of a node might be its *sociability* rather than its *mobility*.

Definition 3.5 (Sociability). Let us look at a node i during a time interval $t_n = [(n-1)T, nT]$, where T is the duration of the interval. Let us further define the indicator function $I_{ij}(t)$, which shows whether nodes i and j are neighbors at time t (e.g., have performed an “association” with Bluetooth). Finally, let $N_i(n)$ denote the set of nodes j such that

$$N_i(n) = \{j \neq i : \exists t \in [(n-1)T, nT] \text{ for which } I_{ij}(t) = 1\}.$$

Then, we define the sociability $S_i(n)$ of node i during the interval t_n as $S_i(n) = \frac{\|N_i(n)\|}{T}$.

In general, the sociability value of a node will be a function of the time interval during which it is measured. If a node’s statistical behavior varies over time, then its sociability index might also change between time intervals. This implies that it might be more appropriate for a node to maintain a *running average* of its *perceived* sociability index, rather than just looking at the previous interval.

Definition 3.6 (MSF Spraying). Each node i maintains a running average for the “sociability index” \hat{S}_i as follows: for a given time interval $t_n = [(n-1)T, nT]$ (T = sliding window duration), it counts the number of unique node IDs encountered, $N_i(n)$.⁴ Then, at the end of window n , it updates \hat{S}_i as

$$\hat{S}_i = (1 - \alpha)\hat{S}_i + \alpha \frac{N_i(n)}{T},$$

and proceeds to the next interval t_{n+1} . α is the weight given to the current measurement in the sliding window mechanism. Finally, the utility of node i is given by $U_i(j) = U_i = \hat{S}_i, \forall j$. This is another example of a DI utility function.

4 SIMULATION RESULTS

We have described a number of utility-based replication protocols that aim at improving the performance of greedy replication schemes. In this section, we are going to use simulations to evaluate all the different flavors we have discussed, namely, LSF, MMF, and MSF, against the basic (greedy) replication scheme (i.e., tree-based Spray and Wait [16], [17]). We have used a custom discrete event-driven simulator using both synthetic mobility as well as some real mobility traces from the DieselNet project [53], [54] to evaluate and compare the performance of the different algorithms. A simplified version of the slotted Carrier-Sense Multiple Access (CSMA) MAC protocol has been implemented. The simulator code and details about the various scenarios used (network parameters, traffic load, and protocol parameters) are publicly available in [55].

For all synthetic mobility scenarios, we assume that 100 nodes move according to the “Community-based




Fig. 1. Scenario 1: performance improvement of LSF Spraying over greedy Spray and Wait (SW) for different configurations; K = node’s transmission range, L = number of copies.

Mobility Model” (CBM) [56], which is motivated by real mobility traces like the ones found in [54]. In the CBM model, each node has its own small community (network size = 500×500 , community size = 50×50) inside which it moves preferentially for the majority of time (e.g., the user’s department building on a campus). Every now and then, it leaves its community (with probability $1 - p_l$), roams around the network for sometime (e.g., going to the cafeteria, library), and then decides to return to its community (with probability $1 - p_r$). Finally, to capture node heterogeneity, each node may have different mobility characteristics (p_l, p_r) in addition to different communities. Node speed is always one network unit/time unit, when a node is moving, and thus, differentiation between node class behavior is achieved only by manipulating the above transition probabilities and community pause times. The work in [57], where a somewhat “richer” version of this model (multitiered communities and time-dependent behavior) has been used, shows that the CBM model closely matches existing traces, confirming that this model better describes real mobility patterns. Finally, we have used here a simple traffic model, with a number of messages (400 unless otherwise stated) routed between uniformly chosen nodes. We intend to look into more sophisticated traffic models in the future.

4.1 LSF Spraying

The first scenario considered (*Scenario 1*) consists of *four different types of nodes* (e.g., to capture a hybrid metropolitan network): (*local nodes*) 40 percent of the nodes move locally most of the time ($p_l \in [0.85, 0.95]$) but may occasionally roam in the whole network ($p_r \in [0.1, 0.2]$); (*community nodes*) 40 percent of the nodes move only inside their own community ($p_l = 1, p_r = 0$); (*roaming nodes*) 10 percent of the nodes roam quite often outside their community ($p_l \in [0.7, 0.8], p_r \in [0.3, 0.5]$); and (*base stations*) 10 percent of the nodes are static and uniformly distributed in the network, corresponding, for example, to base stations or static sensors.

In Fig. 1, we depict the improvement in delivery delay, compared to that of greedy Spray and Wait (“SW”), when message replicas are handed over using LSF Spraying (rule 1 is used—threshold values U_{th} for this scenario were 500–1,000 time units). As can be seen there, taking into account the age of last encounter when spraying the L copies can indeed improve performance.

To further test the benefits of doing utility-based replication, we compare the performance of LSF spraying

4. If T is not too large, the overhead of maintaining a list of unique node IDs and lookup time when a new node is found is not significant.

TABLE 1
Performance Comparison for DieselNet Traces (Scenario 2)

protocol	transmissions (<i>protocol</i> <i>epidemic</i>)	delivery ratio	end-to-end delay (<i>protocol</i> <i>epidemic</i>)
L = 2 copies			
greedy	0.074	0.489	1.59
LSF (cold)	0.058	0.63 (29%)	1.44 (10%)
LSF (warm)	0.084	0.844 (72%)	1.29 (22%)
L = 4 copies			
greedy	0.13	0.676	1.41
LSF (cold)	0.076	0.695 (3%)	1.355 (4%)
LSF (warm)	0.117	0.841 (24%)	1.25 (13%)

against greedy spraying and epidemic routing in a trace-based scenario (*Scenario 2*). Specifically, we utilize traces from the DieselNet project (available at the CRAWDAD repository [54]), which record the encounters between buses of a public transportation network connecting various campuses in Amherst, in an area of approximately 150 square miles. We have performed simulations for a number of different traces from this repository, observing similar behaviors. These traces consist of only a small number of nodes (up to 24 buses), and the improvement achieved is naturally expected to be smaller than the one observed for the previous scenario of 100 nodes. Nevertheless, even in this small scenario, the benefits of utility-based replication are clear.

Table 1 shows results corresponding to trace #2,182,005. “Transmissions” shows the reduction in transmissions compared to epidemic. “Delivery Ratio” is the percentage of messages delivered. End-to-End Delay is the increase in delay compared to epidemic routing. We choose to compare our results against epidemic routing for a light traffic load, for which epidemic is close to optimal in terms of delay and delivery ratio. The values in the parentheses next to the delivery ratio and delay values for LSF correspond to the percentage improvement over greedy.

“LSF (cold)” means that LSF is run directly on the trace for that day ($U_{th} = 1,000$ -5,000 time units) without any previously collected history of encounters from past days. When two copies are used per message, you can see that LSF has a 29 percent higher delivery ratio and a 10 percent smaller delivery delay, compared to the greedy algorithm. However, in a real life situation where the protocol will have been in operation for days or weeks, we expect past encounters from previous days to serve as valuable hints when choosing a relay. Unfortunately, the available traces are collected over a large period, and it is not clear if any of these and which ones correspond to consecutive days. Therefore, to “recreate” a scenario with consecutive days, we assume that the same trace is repeated for two consecutive days, and we examine the potential improvement by LSF on the second day (we denote this protocol as “LSF (warm)”).⁵ In the “warm” version, nodes are run normally in the first day, collecting encounter statistics, but without any traffic generated; traffic is generated and performance is measured only during the second day. Both the cold and warm versions are threshold-based (rule 1);




Fig. 2. Scenario 1: performance improvement of MMF spraying over greedy SW for two different connectivity scenarios (sparse and almost connected); L is the number of copies.

however, in the warm version, U_{th} is made large enough (58,000 time units, which corresponds to more than one day’s time in the trace) to ensure that encounter statistics collected from day 1, carry over to day 2.

“LSF (warm)” shows considerably higher improvement (72 percent higher delivery ratio and 22 percent better delays, for $L = 2$), demonstrating the power of utility-based replication, even in this small scenario. Finally, as the percentage of total nodes that get a copy increases ($L = 4$, i.e., around 20 percent of all nodes get a copy of a message), mistakes made by greedy replication are not as costly, and the improvement achieved by LSF is smaller. Summarizing, these trace-based results confirm the trends observed in the larger synthetic traces.

4.2 MMF Spraying

Here, we turn our attention to the MMF Spraying scheme. As we mentioned in Section 3, MMF Spraying applies to scenarios where there exist multiple classes of nodes. Unfortunately, there do not exist such traces of hybrid networks, as most of them correspond to only one class of nodes (e.g., just vehicles [53], just pedestrians [35], etc.). For this reason, for MSF and MMF, we have only used synthetic simulations to evaluate their performance.⁶

In Fig. 2, we compare the delivery delay of MMF Spraying and Greedy Spraying in the same scenario, as shown in Section 4.1 (Scenario 1). We compare two levels of knowledge for the MMF scheme: 1) messages are given only to nodes of type “local” or type “roaming” but never to “static” or “community” (denoted as MMF1) and 2) now a relay must also have $\{p_l < 0.9 \text{ AND } p_r > 0.15\}$ in addition to label “local” or “roaming” (MMF2). We show plots for both a very sparse network (< 8 percent of nodes are connected) and a denser one (“flakynet”—around 50 percent of nodes are connected). As can be seen by these plots, by choosing only nodes that actually have a chance to deliver a message, we can get up to 2.5-3 \times reduction in delay in this scenario. Greedy spraying wastes many copies by handing them over to community and static nodes.

Intuitively, the larger the amount of node heterogeneity, the higher the expected improvement when using “smarter” copy replication. In this second scenario, we assume only two types of nodes: 1) p percent are “roaming” nodes that are useful as message relays and 2) $(1 - p)$ percent are “community” nodes that are useless as message carriers, unless the destination lies inside their community (small

5. If a bus is used on different routes from one day to another, one could presumably use the route and time information of the given bus rather than its physical ID, as the encounter related quantity of interest.

6. It would be of interest to try to combine some of the existing traces into a hybrid scenario. However, we defer looking into how and if such an endeavor would be successful for future work.




Fig. 3. Scenario 2: performance improvement of MMF Spraying over greedy Spray and Wait, as a function of the percentage p of "useful" relay nodes, different configurations for K = nodes' transmission range, and L = # of copies.

probability). In Fig. 3, we compare the performance of greedy Spray and Wait and MMF Spraying, which gives copies only to roaming nodes, as the percentage p of roaming nodes decreases.

From this figure, it is evident that the smaller the percentage of useful ("roaming") nodes, the higher the delay improvement by using smarter spraying. This is because greedy forwarding makes an increasing number of errors, wasting more and more copies $((1-p)L)$. On the other hand, we see that the achievable improvement peaks at a given value of p and starts becoming smaller again. At this point, the number of useful relays is so small (smaller than the number of copies) that the smarter Spraying scheme takes a very long time to find them.

4.3 MSF Spraying

In this last scenario, we will examine the more practical scheme, MSF Spraying. Unlike the MMF scheme that uses preprocessed information (LABELs) to identify globally useful relays, MSF tries to estimate this utility online by observing the type and number of encounters the node gets involved into. Fig. 4 compares the delivery delay for a scenario similar to that of Fig. 3, where only a percentage p of nodes is useful as message relays. It compares the performance of MSF to Greedy Spraying, as well as MMF spraying (which knows beforehand which are the useful nodes). The threshold-based MSF version is used (rule 1) with a sociability (N_{th}) threshold chosen empirically between five and ten.⁷ As can be seen there, both MSF and MMF improve the performance of Greedy spraying. What is more, MSF can approximate the performance of MMF without the fore-knowledge that the latter requires. (As a final note, although we do not have space to depict plots for the behavior of the running sociability estimates, we have observed that it manages to identify the "good" nodes after only a few windows.)

5 PROTOCOL TUNING

Results in the previous section demonstrate that utility-based replication protocols can achieve significantly superior performance, when compared to greedy replication, in scenarios where node capabilities and behavior are heterogeneous. In this section, we discuss how to tune the main parameters of the proposed protocols, in practice, in order to achieve such good performance. And, in addressing




Fig. 4. Performance comparison of MSF, MMF, and greedy Spray and Wait, as a function of the percentage p of "useful" relay nodes.

some practical issues, we also describe possible complementary mechanisms to the proposed protocols.

5.1 Sliding Window for MSF Sociability Estimation

The MSF replication algorithm is the most generic and practical, among the ones presented, as it can identify "social" nodes *online*. As was shown, it can achieve as good a performance as algorithms that know the "good" relays in advance (e.g., MMF). However, setting the sliding window parameters in MSF correctly (window duration T and α) is necessary to achieve this.

Let us define the stochastic sequence $\{N_i(n), n = 1, 2, \dots\}$, where $N_i(n)$ is the number of unique node IDs encountered by node i during window $t_n : [(n-1)T, nT]$. Whether or not the running sociability estimate \hat{S}_n , based on *past* measurements in t_{n-1}, t_{n-2}, \dots , will be a useful predictor of *future* interactions (i.e., $N_i(n+1), N_i(n+2), \dots$) depends on the amount of "structure" in the interaction and mobility patterns of the nodes involved, the window parameters, and the desired "horizon" for the prediction.

Stationary Case. If the mobility process of all nodes in the network is *stationary*, then the sequence $\{N_i\}$ will also be stationary. Thus, $E[N_i(n)] = E[N_i], \forall n$ and similarly

$$E[S_i(t_n)] = E[S_i(T)], \forall n.$$

The same holds for higher moments of the sequence (i.e., $E[(N_i(n))^k] = E[(N_i)^k], \forall n, k$). This implies that it is not important during which time window one looks at a node's behavior. A node with a highly social past behavior will also have a social future behavior. An example of this could be if all nodes move according to the Random Direction model, but different nodes move with different average speeds and/or pause times. A node i with higher speed and shorter pauses than another node j will meet, *on average*, more new nodes in the same amount of time than j at any time scale.⁸

Since the statistics of the process do not change over time in the stationary case, what matters is choosing the value of α so that our sliding window estimator works efficiently. This depends on the variance $\sigma_{N_i} = E[(N_i)^2]$ of the number of nodes encountered in each window. If the sociability variance of a node during different windows is high, it means that a small α value is needed to smoothen the estimate and ensure convergence. However, the smaller the α , the slower this convergence. One solution is for each

8. We assume here that no encounters are missed. In reality, if a node moves very fast, it might not discover some contact opportunities which are short-lived, due to the workings of the neighbor discovery mechanism [35].

7. We have used $L = 10$ copies per message for all schemes, Tx range $K = 30$, the sliding window for MSF is 1,000 time units, and $\alpha = 0.8$.

node to maintain also the variance of the sampled sociability in addition to the mean. This incurs little overhead and could be used to tune α according to the measured variance. Another solution would be to start with a relatively large α (aggressive) and keep increasing it until the oscillations of the running average from window to window, $E[\hat{S}_i(t_n) - \hat{S}_i(t_{n-1})]$, are smaller than some desirable threshold ϵ .

Cyclo-Stationary (Periodic) Case. In real life, people (and sometimes, animals) often tend to have different mobility/interaction patterns during different times of the day. It has been frequently observed that real-life mobility exhibits such time periodicity [58]. In this case, the statistics of the encounter sequence N_i might change as a function of time-of-day. This implies that for some scenarios, N_i (and thus, S_i) could be modeled as a *cyclo-stationary* process.

In these processes, if the sliding window T duration is smaller than the process' cycle, then what happened in the previous time window cannot necessarily predict what will happen in the next. Yet, stationarity may also emerge here given a large enough time horizon for the prediction. If the duration of such a period is not prohibitive compared to the desired delivery delay,⁹ setting the sliding window T to this value would allow one to reliably predict future delivery probability in the next cycle.

On the other hand, if finer granularity is needed for the prediction (i.e., targeted delays are much smaller than this period), one is forced to use a shorter sliding window. In this case, one could maintain multiple sociability values, corresponding to different time periods of the day (e.g., "morning," 2-4 p.m., etc.): then, each sample measured will affect only the time bin to which it corresponds, and inquiries about future delivery probability will take into account the sociability estimate corresponding to the time period(s) included in the horizon until the message in hand expires.

Nonstationary Case. In some cases, the mobility process of the nodes might be too complex to be modeled as a stationary or cyclo-stationary process. Nevertheless, mobility/interaction behavior in real life is never absolutely random. Even if the statistics (long-term behavior) of the process are changing, one would assume that these do so slowly. (Note that we do not here talk about the event of measuring different sociability values during different windows, something expected and predicted by the variance of the process, but rather the random process parameters changing.) In this case, the running average should still be able to follow and adapt to these changes, even so with a small lag that might just introduce a bit of error in the predicted delivery probability.

5.2 Number of Copies

Another important protocol parameter is the number of copies L to be used. This number has been calculated for Greedy Spraying in a homogeneous network, in order to achieve a desired *transmission-vs-delay* trade-off [16]. The same procedure could be modified for the case of MMF spraying. The main idea is that if there are M total nodes and MMF uses only p percent of them based on their labels,

then this is equivalent to a network of pM nodes (some care is needed though, regarding the effect of the other $(1-p)M$ nodes in the minimum/optimal delay [16]).

In the case of MSF Spraying, a different approach could be taken, as described by Lemma 5.1.

Lemma 5.1. Assume that there are M total nodes in a network moving independently of each other, and messages have a TTL value of T_{max} . Let further each node use the MSF algorithm (rule 1) with sliding window $T = T_{max}$ and threshold sociability value $S_{th} = \frac{N_{th}}{T_{max}}$. If the desired delivery probability for each message is R_{min} , then the number of copies per message L_{min} and sociability threshold must satisfy

$$1 - \left(1 - \frac{N_{th}}{N}\right)^{L_{min}} \geq R_{min}, \quad (3)$$

where N is the total number of nodes.

The above lemma can be derived by setting the sliding window duration T equal to the TTL value (T_{max}). Then, the sociability threshold would become $\frac{N_{th}}{T_{max}}$. N_{th} becomes an estimate of the unique number of nodes encountered within the message's lifetime and thus of the delivery probability $\frac{N_{th}}{N}$ of one relay, assuming that relays encounter other nodes in the network independently. If L_{min} relays are used in total, the probability of delivery before the TTL expires is thus $1 - \left(1 - \frac{N_{th}}{N}\right)^{L_{min}}$. Although this might not be realistic in practice, it could serve to either find a minimum S_{th} given L_{min} or L_{min} given S_{th} , if the sliding window is $T = T_{max}$. An interesting question arises when one considers how to scale this value for a smaller time window. For example, consider how the parameters used for MSF in the case of Fig. 4 fit this equation. The main problem in this case is that the sliding window there is $T = 1,000$ (with $N_{th} = 5-10$), while $T_{max} = 10,000$; thus, we have only an estimate of the average number of nodes encountered during $\frac{1}{10}$ of a packet's lifetime. If $N_{1,000}$ is the average sociability when the window is $T = 1,000$, what can we say about $N_{10,000}$? Linear scaling (i.e., $N_{th} = 50-100$) gives too optimistic results (0.99 delivery ratio). The correct value of N_{th} corresponding to a sliding window of $T_{max} = 10,000$, to achieve $R_{min} \geq 0.9$ (as in Fig. 4) seems to be around $N_{th} = 25$. We believe that prediction methods like the one used in [16] could provide the answer to this scaling question, but we defer this for future work.

In general, the joint problem of choosing both the number of copies and the copy bearers optimally is difficult. We assume here that either the number of copies is set by the application (e.g., credits per message) or the protocol is fixed and the network parameters known. Nevertheless, we next discuss a methodology that could potentially address the joint optimization problem.

5.3 Copy Allocation as Portfolio Management

Whether the nodes' mobility process is stationary or not, maintaining higher moments of a node's past sociability can allow a node to make more "educated" decisions when allocating message copies. For example, assume that each node maintains a tuple $\{\hat{S}_i, \hat{\sigma}_{S_i}\}$, that is, estimates of the node's average sociability and its variance, respectively. Let, for example, a node carrying an extra copy of a message encounter nodes i and j . Let further $\hat{S}_i > \hat{S}_j$ but $\hat{\sigma}_{S_i} \gg \hat{\sigma}_{S_j}$.

9. Given the delay-tolerant nature of the targeted services (e.g., advertisements, e-mail, message boards) and some of the targeted environments (e.g., developing countries, remote villages, etc.), it is not difficult to imagine that for some users, delays as large as one day or more (i.e., equal to the daily process' cycle) could be acceptable.

This means that although node i is more sociable, in general, than j , it is also much more “risky”: its future behavior cannot be reliably predicted. Also, whether i or j should get the extra copy should also depend on who the other relays having received a copy are. A list of the relays that have received a copy can be maintained, with a predicted delivery profile (expected delay and variance) based on their joint characteristics. Then, *every new copy forwarded should aim to improve this profile*. For example, if we know that a number of “consistently” social nodes (i.e., good sociability with small variance) already have a copy, it might be worthwhile to risk giving a copy to i , which might have some chance of delivering the copy faster than other relays, as opposed to using yet another “average” relay. This results into more complex, but also more powerful forwarding decisions, as opposed to simply comparing average values.

Although we have described this procedure in the context of MSF, it is important to note that it could also be applied to other utility-based replication variants, e.g., LSF. In the case of LSF, a tuple $\{\hat{T}_{id}, \hat{\sigma}_{T_{id}}\}$ with estimates of the average and variance of the intermeeting time would be maintained. In general, the problem of allocating copies to nodes with different expected “return” (e.g., delivery delay) and different “risk” (e.g., variance of predicted delivery delay), is, in some aspects, equivalent to optimal asset allocation and portfolio theory in finance [59]. An interesting effort uses ideas from modern portfolio theory to address some reliability issues [38]. We intend to look more into this issue in future work.

5.4 Avoiding Node Overloading

If message replicas are handed over to a few good nodes, then these nodes might absorb a disproportionate amount of workload compared to the average node. Accepting a copy implies a particular overhead for a node in terms of buffer space, battery power, and bandwidth consumption. Nevertheless, in many scenarios of interest node, overloading is not a concern. For example, in situations where some nodes (e.g., buses with no energy constraints [53]) have considerably more resources than the average node (e.g., a pedestrian carrying a PDA), it is desirable to direct more traffic to the overprovisioned nodes. Furthermore, if overusing a particular node resource is a concern, one could incorporate the usage of this particular resource into a (hybrid) utility function that aims at optimizing a combination of different goals. For example:

Remaining battery level. A node may use a hybrid utility function $\hat{U} = f(U_{orig}, power)$, where U_{orig} captures the probability of message delivery (i.e., $U_{orig} = \{U^{LSF}, U^{MMF}, U^{MSF}\}$) and $power$ is the remaining battery level (taking values in $[0, 1]$). Then, the forwarding rule could be defined as:

- If $power > power_{min}$, then $\hat{U} = U_{orig}$.
- If $power < power_{min}$, then $\hat{U} = 0$.

This policy sets a “cutoff” ($power_{min}$) that prevents nodes to serve as relays when their remaining lifetime is below a certain level. Another policy could weigh the advertised utility by the remaining power:

$$\hat{U} = power \times U_{orig} \quad (\text{linearly}),$$

$$\hat{U} = 1 - \frac{1}{(U_{orig})^{power}} \quad (\text{nonlinearly}).$$

A similar mechanism could be employed to prevent a node from serving as a relay if its buffer occupancy drops below a threshold. Thus, under light loads, the protocol operates with routing efficiency in mind, while at high loads, it makes load-balancing a factor.

5.5 Security Considerations: Counterfeiting Utility

Allowing “gravitational points” that attract a high proportion of the network load can also become a security “loophole.” Ideally, the node that forwards the copy(ies) trusts the next hop to be stating its true, measured utility. Thus, *if a malicious node decides to advertise a very high (fake) utility value, it can easily absorb a large number of messages*. It could then just drop them (DoS) or try to read/decode them (compromise privacy).

The proposal of security solutions for delay-tolerant networking is beyond the scope of this paper. We should point out that it is an area of active research. For instance, some security-related issues arising from intermittent connectivity have been partly addressed in [60]. Reputation mechanisms based on monitoring neighboring transmission are revisited in [61] since such monitoring is difficult in sparse networks. The studies reported in [21], [62] show that the very nature of epidemic replication makes it quite resilient to many types of attacks. Finally, it has been shown that mobility itself could facilitate the establishment of a robust trust system on the fly [63].

6 PERFORMANCE ANALYSIS OF ENCOUNTER-BASED ROUTING IN HETEROGENEOUS NETWORKS

Having concentrated in the previous sections on the protocol performance and tuning, in this last section, we embark on developing appropriate theoretical models for encounter-based message dissemination in heterogeneous networks. These models not only aim to help us analyze the performance of the routing schemes presented, but can also serve as the foundation to analyze various other encounter-based processes (e.g., wireless virus spread [64]) in networks consisting of mixed populations of nodes.

Our theoretical framework is based on an appropriate *fluid model* formulation for epidemic routing [19], *in a heterogeneous setting*. The basic concepts and terminology are borrowed from the field of Epidemiology and Mathematical Biology [26]. Nodes having already received a copy (we call these nodes “Infected” I) may forward an additional copy to a node not having yet received one (we call this node “Susceptible” S). We are then usually interested in the *expected time* until a given node (e.g., destination) gets infected (i.e., receives the message) or the probability of this occurrence (delivery ratio). A fluid model can then be formulated that captures the rate of message propagation (“infection”) among nodes and can be used to calculate the above quantities. Our model extends the fluid model formulation of [27], which was proposed for homogeneous networks with a single class of nodes. Although fluid models are only a *deterministic approximation* of the *stochastic* spreading of messages using epidemic

TABLE 2
Notation

N	size of network area
M	total number of nodes
M_i	number of nodes of class i
β_{ij}	mixing rate between nodes of classes i and j
$EM_{ij} = \frac{1}{\beta_{ij}}$	exp. meet time between nodes of class i and j
$I_i(t)$	infected nodes of class i at time t
$D_{protocol}$	delivery delay of protocol "protocol"
$P(t)$	complementary cumulative distribution function of random variable $X(t)$: $P(t) = Prob(X(t) > t)$

routing, they produce accurate results for moderate to high numbers of networks nodes [27].

We have chosen to start by modeling epidemic routing first for the following reasons: 1) epidemic routing [19] is the first, and among the most generic opportunistic routing protocols proposed with many variants having followed (e.g., Prophet [20], SLEF [62], etc.); 2) its delay is *minimum* under ideal conditions (little or no contention for shared resources like bandwidth and buffer space), and thus can serve as an optimality baseline; and 3) we will use it to describe a version of controlled replication (greedy and utility-based), which is equivalent to the one originally presented [22], but is more amenable to our fluid-model-based analysis. Having developed this basic model for epidemic routing, we then use it to quantify the performance of Greedy Replication when node mobility is heterogenous (the homogeneous case is treated in [17], [22]). Finally, we analyze how good a delay we can achieve, as compared to the optimal case, by using Smart Replication algorithms like MSF or MMF.

Before we proceed, it is important to define node heterogeneity in a rigorous manner. *Our model for a heterogeneous network consists of multiple classes of nodes (k total), with different intra and interclass meeting characteristics.* The reason for this choice is twofold: 1) this setting can successfully capture two scenarios of high interest: first, scenarios where nodes are naturally separated into communities or groups with high intragroup meeting rate (e.g., same department, affiliation, and family), and second, scenarios where some subsets/classes of nodes have different characteristics or capabilities than others (e.g., vehicles, pedestrians, base stations, etc.) and 2) this multi-population model is more amenable to fluid model analysis; it has also been used often in epidemiology from which we might be able to borrow some results or formalisms.

In the remainder of this section, we will assume the following:

1. There are M nodes in total, each moving according to some *stochastic* mobility model in a finite area. This mobility is neither known in advance [2] nor enforced [29].
2. Each of the M nodes belongs to a particular mobility class i , among k possible ones. Nodes in different classes may move according to different mobility models (e.g., some nodes perform random walks, while others perform random direction mobility) or simply have different mobility parameters (e.g., all nodes perform random waypoint mobility, but nodes in one class have high average speed and

small pause times). The total number of nodes in class i is denoted as M_i with $\sum_{i=1}^k M_i = M$.

3. Let a node i belong to mobility class $i \in [1, k]$ and another node j belong to mobility class $j \in [1, k]$. Further, let the two nodes start from their stationary distributions at time 0. The *meeting time* (M_{ij}) between them is defined as the time it takes to first come within range of each other, that is, $M_{ij} = \min_t \{t : \|X_i(t) - X_j(t)\| \leq K\}$, where K is the maximum transmission range and $X_i(t)$ is the position of node i at time t .
4. We define the meeting or "mixing" rate between the two classes i and j , β_{ij} as the inverse of the expected meeting time for these classes: $\beta_{ij} = \frac{1}{EM_{ij}}$. The mixing rate is a measure of how often a node of one class meets a given other node of another (or the same) class. We will assume that all meeting times M_{ij} , $i, j \in [1, k]$, are approximately exponentially distributed (or have an exponential tail [48]) with expected value $EM_{ij} = \frac{1}{\beta_{ij}}$. It has been shown that a number of popular mobility models have such exponential tails (e.g., Random Walk, Random Waypoint, Random Direction, and Community-based Mobility [41], [56]).
5. $I_i(t)$ denotes the number of nodes of type i that have already been infected (i.e., have received a copy) at time t .

6.1 Epidemic Routing

In this section, we start by modeling the performance of the Epidemic Routing protocol [19] in heterogenous networks. A similar formulation for multiple populations can be found in the context of Mathematical Biology, yet with different quantities of interest in mind (e.g., epidemic threshold) [45].

Lemma 6.1. *Let us assume that there are k classes of nodes and one message from a node in a class $s \in [1, k]$ originates at time $t = 0$. If this message is propagated using epidemic routing and no "buffer cleaning" policy,¹⁰ the spreading of this message is given by the solution of the following system of rate equations:*

$$\frac{dI_i(t)}{dt} = (M_i - I_i(t)) \sum_{j=1}^k I_j(t) \beta_{ji}(t), \forall i \in [1, k], \quad (4)$$

$$I_s(0) = 1, I_j(0) = 0, \quad \forall j \neq s.$$

Proof. A fluid model approximation for a single class of nodes has been formulated and justified in [27]. Here, we extend this formulation into multiple classes, using similar arguments and assumptions about scaling as [27]. Let us assume that there are $I_i(t)$ nodes of class i that have received a copy of the message (i.e., are "infected") at time t . The number of susceptible nodes of class i is $M_i - I_i(t)$. If there are also $I_j(t)$ nodes infected from class j and *individual* nodes of classes i and j

10. In [17], it is suggested that "antipackets" are propagated back to infected nodes after the message reaches its destination, in order to release some buffer space and avoid further infection. These nodes would then be considered "Recovered" (R), and an SIR model would be more appropriate. Without loss of generality, here we decide to treat the simpler case of no antipackets (i.e., "No-Immune" version in the paper).




Fig. 5. Number of infected nodes in class 1 (I_1 —static nodes) and class 2 (I_2 —mobile nodes) as a function of time. We have normalized time and mixing rates such that $\beta_{11} = 0, \beta_{12} = 1, \beta_{22} = 2$.

encounter each other (i.e., mix) with rate β_{ij} , then the probability of a new node of class i getting infected by a class j node at the next time instant $\Delta t \rightarrow dt$ is

$$I_i(t + dt) - I_i(t) = \beta_{ij}(M_i - I_i(t))I_j(t)dt.$$

Taking also into account all other classes $j \in [1, k]$ of nodes that can (independently) infect nodes of class i gives us the system of (4). Initial conditions simply state that only a node from class s is infected at first. \square

The next theorem derives the complementary cumulative probability distribution function (CCDF) of the delay until a message is delivered to a randomly chosen destination. The CCDF completely describes the delivery delay for epidemic routing (expected value and higher moments).

Theorem 6.1. *Let us assume that the destination for the message propagated in Lemma 6.1 belongs to class “d.” Further, let the CCDF of the delay until the destination receives the message be denoted as $P = P(D_{epid} > t)$. Then, P is given by solving the following equation:*

$$\frac{dP(t)}{dt} = -P(t) \sum_i \beta_{id} I_i(t), \quad (5)$$

where $I_i(t)$ are given by Lemma 6.1.

Proof. Let us calculate the probability $P(D_{epid} > t + dt)$. It is easy to see that

$$P(D_{epid} > t + dt) = P(D_{epid} > t) \times (1 - \text{Prob}(t \leq D_{epid} \leq t + dt)).$$

Since meetings between nodes are approximated by a Poisson process, the probability of the destination being found at any instant of duration dt is equal to the aggregate meeting rate of infected nodes with the destination at time t times dt : $\sum_i I_i(t) \beta_{id} dt$. Thus,

$$\begin{aligned} P(t + dt) &= P(t) \left(1 - \sum_i I_i(t) \beta_{id} dt \right), \\ \Rightarrow \frac{P(t + dt) - P(t)}{dt} &= -P(t) \sum_i I_i(t) \beta_{id}. \end{aligned}$$

\square

Based on the above density function, the expected delay for epidemic routing in this multiclass population follows easily from elementary probability.

Corollary 6.1. *The expected delay of epidemic routing in a network of k node populations is given by $ED_{opt} = \int P(t) dt$.*




Fig. 6. CCDF of the delivery delay for the two scenarios of Fig. 5. Destination is chosen uniformly in class 1.

Two mobility classes ($k = 2$). To get some insight into the above formulation, we will look in more detail into the special case of only two classes of nodes, as in Section 3.1. Specifically, we will assume that the network comprises: (class 1) a fraction of nodes $1 - p$ that move more slowly and rarely than nodes in the second class ($b_{22} \gg b_{11}$) or are static ($b_{11} = 0$), (class 2) a fraction of nodes p that move frequently and/or fast around the network (e.g., vehicles). In this case, the system of ODEs of Lemma 6.1 becomes

$$\frac{dI_1(t)}{dt} = (M_1 - I_1(t))(I_1(t)\beta_{11} + I_2(t)\beta_{12}), \quad (6)$$

$$\frac{dI_2(t)}{dt} = (M_2 - I_2(t))(I_1(t)\beta_{12} + I_2(t)\beta_{22}), \quad (7)$$

$$I_1(0) = c_1, I_2(0) = c_2. \quad (8)$$

In Fig. 5, we depict $I_1(t), I_2(t)$, the total number of infected nodes of each class as a function of time, for different values for M_1, M_2 ; and in Fig. 6, we depict the respective cumulative probability density functions. It is easy to observe from these figures that when there are more “fast” nodes in the network (class 2) infection is faster, and delivery delay smaller.

Finally, in Fig. 7, we compare simulation results for the delay of epidemic routing against analytical values calculated according to Corollary 6.1, again for a scenario of two populations of nodes, as a function of the ratio between nodes of the two classes. We observe that theoretical and simulation results are closely matched, which confirm the accuracy of our fluid model.

6.2 Controlled Replication (Spraying)

Here, we will turn our attention into analyzing the delay of Controlled Replication, both greedy- and utility-based. The




Fig. 7. Comparison of analytical and simulation results for the expected delay of epidemic routing; $M = 300$ total nodes in a $1,000 \times 1,000$ network with $K = 30$; p of the nodes move according to the Random Direction model, while $1 - p$ are static.

delay of greedy replication (Spray and Wait) has been analyzed using probability theory and Markov Chains in [16], [17], [22] for homogeneous node scenarios. To calculate the delay in heterogeneous networks, we will use again fluid models, as in the previous section, due to their power to model population mixing problems. However, instead of using the scheme described in Definition 3.1, we will use the variant of Definition 6.1. This one is more amenable to a fluid model analysis, and has equivalent performance to the original scheme, *on average*.

Definition 6.1 (Time-Bound Spray and Wait). *Let us assume that a message gets created at time T_{start} . Let further T_{spray} be such that $I(T_{spray}) = \sum_i I_i(T_{spray}) = L$, where $I_i(t)$ are given by Lemma 6.1. Then,*

- (spray phase) all nodes (source and relays) perform epidemic forwarding (i.e., give a copy to every new node encountered) until time $T_{start} + T_{spray}$;
- (wait phase) after time $T_{start} + T_{spray}$, any node that has received a message copy, may only give a new copy to the message destination.

The above scheme emulates greedy, tree-based replication with L copies (see Section 3.1).

6.2.1 Greedy Replication

Theorem 6.2. *Let us assume that there are k classes of nodes in a network with M total nodes ($M = \sum_{i=1}^k M_i$). Let further a message originate at time $t = 0$ from a node in a class $s \in [1, k]$ and destined to a node of class $d \in [1, k]$. Let further this message be routed using the time-bound Spray and Wait protocol of Definition 6.1. Then, the expected delivery delay of the message, ED_{sw} , is upper bounded by*

$$ED_{sw} \leq T_{spray} + \frac{1}{M} \frac{M - \sum_{i=1}^k I_i(T_{spray})}{\sum_{i=1}^k \beta_{id} [I_i(T_{spray})]}, \quad (9)$$

where $I_i(T_{spray})$ are given by Lemma 6.1.

Proof. T_{spray} is the time during which message copies are distributed to relays (using epidemic spreading as described in Definition 6.1). During this time $\sum_{i=1}^k I_i(T_{spray})$, copies are given in total, where $I_i(t), i \in [1, k]$ are given by (4). Since source-destination pairs are chosen uniformly among all nodes, the probability that the destination has not been found during the copy distribution phase is equal to

$$P(wait) = \frac{M - \sum_{i=1}^k I_i(T_{spray})}{M}.$$

In that case, each of the L nodes will hold on to its message copy until it encounters the destination itself (or TTL expires). Let us denote this additional time as T_{wait} , and let at the end of T_{spray} exist $[I_i(T_{spray})]$ relays of class i (since no fractions of nodes can exist). If we assume that each one of them meets the destination independently, and the individual meeting times are exponential with average $\frac{1}{\beta_{id}}$ for class i , then

$$T_{wait} = \frac{1}{\sum_{i=1}^k [I_i(T_{spray})] \beta_{id}}.$$

Putting it altogether

$$ED_{sw} \leq T_{spray} + \frac{M - \sum_{i=1}^k I_i(T_{spray})}{M} \frac{1}{\sum_{i=1}^k \beta_{id} [I_i(T_{spray})]}.$$

The reason why the above is an upper bound is that we always count the whole spraying delay T_{spray} , even if the destination was actually found before the end of the replication timer. Nonetheless, if $L \ll M$ (the case we are most interested in, since we want to drastically reduce the number of nodes that are burdened with a copy), then the above bound becomes tight, similar to the single-population case [22]. \square

6.2.2 Utility-Based Replication

Next, we will look into the delay of Utility-based Replication that allows only nodes in particular classes to act as relays. Although this is typically the MMF routing protocol, it can also model MSF routing with appropriately chosen thresholds. We start with a special case that is easier to handle, where only one class of nodes can act as relays. We can adapt the results of [16] to calculate the delay for this case, without resorting to fluid models. We use the fluid model next to address the general case.

Lemma 6.2. *Let us assume that a message, from a node in a class $s \in [1, k]$ to a node of class $d \in [1, k]$, is routed with MMF spraying using only nodes of some class $r \in [1, k]$. Further, let $ED_{MMF}(L+1)$ be the expected delay of the scheme, when $L+1$ copies are used. Finally, let $ED(i)$ be the expected remaining delay after i message copies have been spread. Then, $ED(1) \approx ED_{MMF}(L+1)$, where $ED(1)$ is calculated by the following system of recursive equations:*

$$ED(i) = ET_{next}(i) + P_{next}(i)ED(i+1), \quad i \in \left[1, \frac{L}{2}\right],$$

$$ED(i) = ET_{next}(i) + P_{next}(i) \left[\frac{2i-L}{i} ED(i) + \frac{L-i}{i} ED(i+1) \right], \quad i \in \left[\frac{L}{2}+1, L-1\right],$$

$$ED(L+1) = \frac{1}{L\beta_{rd} + \beta_{sd}},$$

where

$$ET_{next}(i) = \frac{1}{(i-1)[(M_r - i + 1)(\beta_{rr} + \frac{\beta_{sr}}{i-1}) + \beta_{rd}] + \beta_{sd}},$$

$$P_{next}(i) = \frac{(M_r - i + 1)(\beta_{rr}(i-1) + \beta_{sr})}{(M_r - i)(\beta_{rr}(i-1) + \beta_{sr}) + (i-1)\beta_{rd} + \beta_{sd}}.$$

Proof. This proof follows a similar methodology as that in [16, Theorem 4.2]. Let us look into the case, when there are i nodes ($i < L$) that have one or more copies. Since the source *always* has at least one copy, this would imply that there are $i-1$ relays of class r and one of class s . Further, let us assume that among the i nodes with copies, X_i of them have more than one copy ("active") and are allowed thus to forward copies further to other relays. Since all hitting times are independent and exponentially distributed, nodes with a message copy

(i total) encounter: 1) any of the nodes of class r without a copy ($M_r - i$) with rate $(i - 1)(M_r - i)\beta_{rr} + (M_r - i)\beta_{sr}$ or 2) the destination itself with rate $i\beta_{rd} + \beta_{sd}$. This is the time until the next *encounter of interest* occurs (either a candidate relay or the destination), and we shall denote it as $ET_{next}(i)$. Now, if the node encountered is the destination, the message gets delivered. Otherwise (with probability $P_{next}(i) = \frac{(M_r - i + 1)(\beta_{rr}(i - 1) + \beta_{sr})}{(M_r - i)(\beta_{rr}(i - 1) + \beta_{sr}) + (i - 1)\beta_{rd} + \beta_{sd}}$), the algorithm continues, performing one of the following: 1) with probability $\frac{X_i}{i}$, it is one of the “active” nodes that encountered this other node, and therefore, hands it over half its copies; $i + 1$ nodes have copies now, and an expected time $ED(i + 1)$ remains until delivery and 2) with probability $\frac{i - X_i}{i}$, it was one of the other nodes carrying a message copy that encountered a new node. Since these relays forward only their message copy to its destination, nothing happens, and the remaining time is still $ED(i)$. Putting it altogether

$$ED(i) = ET_{next}(i) + P_{next}(i) \left(\frac{i - X_i}{i} ED(i) + \frac{X_i}{i} ED(i + 1) \right).$$

In a manner similar to [16, Theorem 4.2], we can show that a useful approximation is $X_i = i, i \in [1, L/2]$ and $X_i = L - i, i \in [L/2 + 1, L - 1]$.

Finally, if all L copies have been distributed (1 at source and L at nodes of class r) and the destination has not been found yet, then the remaining time is easily found to be $ED(L + 1) = \frac{1}{L\beta_{rd} + \beta_{sd}}$. \square

Theorem 6.3 addresses the general case where utility-based spraying may use more than one classes of nodes as relays. Its proof follows a similar methodology as Lemma 6.1.

Theorem 6.3. *Let a message be routed using the time-bound Spray and Wait protocol of Definition 6.1. However, assume further that only a subset of classes R can act as relays $R \subset [1, k]$ (i.e., MMF/MSF spraying). Then, the expected delivery delay of the message ED_{MMF} is upper bounded by*

$$ED_{MMF} \leq T_{spray} + \frac{1}{M} \frac{M - \sum_{i=1}^k I_i(T_{spray})}{\sum_{i \in R} \lfloor I_i(T_{spray}) \rfloor \beta_{id}}, \quad (10)$$

where $I_i(T_{spray})$ are given by Lemma 6.1.

Proof. The proof goes along the same lines as that of Theorem 6.2, with the difference that the nodes of classes who *cannot* act as relays ($\notin R$) are ignored. Also, note that $L = \sum_{i \in R} I_i(t)$, the total number of copies distributed after T_{spray} , will be smaller here than the value of Theorem 6.2, everything else equal. To make sure that the average number of copies distributed for MMF spraying is equal, on average, as in the greedy replication case, we need to set $T_{spray}(MMF)$ such that:

$$\begin{aligned} T_{spray}(MMF) &= \arg \min_t \left\{ \sum_{i \in R} I_i(t) = L \right\} \\ &\geq T_{spray}(greedy) = \arg \min_t \left\{ \sum_{i=1}^k I_i(t) = L \right\}. \end{aligned}$$

This implies that the first component in the right-hand side of (10), T_{spray} , will be larger in the MMF case than in the greedy case ((9)). The reason is that in the greedy case, after L nodes with no copy are encountered, the spraying phase always ends; in the utility-based case, the first L nodes do not necessarily become relays (if they do not have high enough utilities). Thus, the spraying phase continues and finishes only after enough nodes have been met (say $L' \geq L$), such that L of them had a high enough utility to receive a copy. For the same reason, the probability that the destination is found during the “spraying” phase for MMF is $P_{util}(wait) = \frac{L'}{M}$; if one of the L' nodes encountered is the destination, the message gets delivered. Since, in the greedy case, this probability is $P_{greedy}(wait) = \frac{L}{M}$, and $L' \geq L$, it follows that the probability of finding the destination in the spraying phase is higher for MMF. In fact, if the subset of candidate relays (classes in R) becomes sufficiently narrow (total number of nodes in classes in R is smaller than L), then the destination is always found during the spraying phase. \square

The following Lemma and its corollary (Corollary 6.2) compare the performance of greedy replication and utility-based replication in multiple-population scenarios:

Lemma 6.3. *The expected remaining delay (“wait phase delay”) after all copies, i.e., L , have been distributed is always smaller in the case of MMF than in the case of greedy replication.*

Proof. Let us assume that $T_{spray}(MMF)$ and $T_{spray}(grd)$ are such that L copies are distributed, on average, for both MMF and greedy replication. Then

$$\begin{aligned} &\sum_{i \in R} I_i^{MMF}(t) \beta_{id} \\ &\stackrel{(1)}{=} \sum_{i \in R} I_i^{grd}(t) \beta_{id} + \sum_{i \in R} (I_i^{MMF}(t) - I_i^{grd}(t)) \beta_{id} \\ &\stackrel{(2)}{\geq} \sum_{i \in R} I_i^{grd}(t) \beta_{id} + \min_{i \in R} \{\beta_{id}\} \sum_{i \in R} (I_i^{MMF}(t) - I_i^{grd}(t)) \\ &\stackrel{(3)}{=} \sum_{i \in R} I_i^{grd}(t) \beta_{id} + \min_{i \in R} \{\beta_{id}\} \sum_{j \in \hat{R}} I_j^{grd}(t) \\ &\stackrel{(4)}{\geq} \sum_{i \in R} I_i^{grd}(t) \beta_{id} + \sum_{j \in \hat{R}} I_j^{grd}(t) \beta_{jd} \\ &\stackrel{(5)}{=} \sum_{j=1}^k I_j^{grd}(t) \beta_{jd}. \end{aligned}$$

Step (1) is simply adding and subtracting $\sum_{i \in R} I_i^{grd}(t) \beta_{id}$. Step (2) follows by replacing each $\beta_{id}, i \in R$ with the minimum value among them. Since the number of copies spread for both the greedy and the utility-based cases is equal to L , it holds that $(\sum_{j \in \hat{R}} I_j^{grd}(t) = I_i^{MMF}(t) - I_i^{grd}(t))$. This simply says that the rest of copies, which in the MMF case are given to some nodes of classes in R , are given to nodes in classes outside R (\hat{R}) in greedy replication. This gives step (3). Step (4) follows from the fact that nodes in any class in R mix faster with the destination than nodes in any class outside R . Thus, $\min_{i \in R} \{\beta_{id}\} \geq \beta_{jd}, \forall j \notin R$. Step (5) is simply merging the two sums.

Consequently,

$$\frac{1}{\sum_{i \in R} [I_i(T_{\text{spray}}(\text{MMF}))] \beta_{id}} \leq \frac{1}{\sum_{i=1}^k [I_i(T_{\text{spray}}(\text{grad}))] \beta_{id}}.$$

□

Corollary 6.2. *Let s and d denote the node class of the source and the destination of a message, respectively, and let $s, d \notin R$, where R is the set of classes that can act as relays in the MMF replication algorithm. Finally, let $\frac{\beta_{kl}}{L\beta_{ij}} > 1, \forall i, j, k \in \bar{R}$, and $\forall l \in R$. Then, the expected delivery delay of the MMF algorithm ED_{MMF} is always smaller than the expected delay of greedy replication ED_{grad} .*

Proof. This can be derived using some calculations from Lemma 6.3 and Theorem 6.3. □

Lemma 6.3 and Corollary 6.2 essentially say that if data sources and sinks do not generally mix but rely on a set of “special,” mobile nodes to carry traffic around, then utility-based replication is *always* faster than greedy replication. This would be the case, for example, when nodes are either static (sensors) or slowly moving (pedestrian) while routers are vehicles like buses [53].

As a final note, we have also compared simulation and analytical results (bounds) for greedy and utility-based replication for heterogeneous networks consisting of two populations of nodes, as in the epidemic routing case (Section 4.1). We have observed that respective results present a similarly good match, so we choose to not include the plots due to space limitations.

7 CONCLUSION

In this work, we have looked into the issue of efficient mobility-assisted or opportunistic routing for networks comprising heterogeneous node populations. We have explored the issue of allocating a fixed budget of message replicas among a number of candidate relays with different capabilities and mobility patterns. To this end, we have proposed a number of simple heuristics that try to capture how probable it is for a given relay to encounter the destination in the near future. Using both theory and simulations, we show that *choosing* the relays carefully rather than picking the first few relays encountered [16], [17], [18] can significantly improve performance in heterogeneous scenarios.

Possible directions for future work include exploring the design of more sophisticated or hybrid heuristics, as well as addressing the issue of *optimally* solving the joint number-of-copies/copy-allocation problem. Yet, we believe that this work clearly shows that utility-based replication presents a desirable and often necessary routing mechanism, when the number of total transmissions needs to be controlled, while at the same time, not all nodes in the network are equally useful for routing. Finally, using utility-based replication rather than greedy replication during the initial distribution phase of message replicas could also have an improvement on the performance of schemes that allow multihop forwarding [16].

ACKNOWLEDGMENTS

This work was performed while all of the authors were at INRIA, Sophia-Antipolis, France.

REFERENCES

- [1] A. Seth, D. Kroeker, M. Zaharia, S. Guo, and S. Keshav, “Low-Cost Communication for Rural Internet Kiosks Using Mechanical Backhaul,” *Proc. ACM MobiCom*, 2006.
- [2] S. Jain, K. Fall, and R. Patra, “Routing in a Delay Tolerant Network,” *Proc. ACM SIGCOMM*, Aug. 2004.
- [3] A. Pentland, R. Fletcher, and A. Hasson, “DakNet: Rethinking Connectivity in Developing Nations,” *Computer*, vol. 37, no. 1, pp. 78–83, Jan. 2004.
- [4] P. Basu and T. Little, “Networked Parking Spaces: Architecture and Applications,” *Proc. IEEE Vehicular Technology Conf.*, 2002.
- [5] M. Papadopoulou and H. Schulzrinne, “Seven Degrees of Separation in Mobile Ad Hoc Networks,” *Proc. IEEE Global Comm. Conf. (GLOBECOM '00)*, 2000.
- [6] M. Motani, V. Srinivasan, and P.S. Nuggehalli, “PeopleNet: Engineering a Wireless Virtual Social Network,” *Proc. ACM MobiCom*, 2005.
- [7] P. Hui, A. Chaintreau, J. Scott, R. Gass, J. Crowcroft, and C. Diot, “Pocket Switched Networks and Human Mobility in Conference Environments,” *Proc. ACM SIGCOMM Workshop Delay Tolerant Networking (WDTN '05)*, 2005.
- [8] Disruption Tolerant Networking, <http://www.darpa.mil/ato/solicit/DTN/>, 2009.
- [9] J. Heidemann, W. Ye, J. Wills, A. Syed, and Y. Li, “Research Challenges and Applications for Underwater Sensor Networking,” *Proc. IEEE Wireless Comm. and Networking Conf.*, 2006.
- [10] Delay Tolerant Networking Research Group, <http://www.dtnrg.org>, 2009.
- [11] E. Brewer, M. Demmer, B. Du, M. Ho, M. Kam, S. Nadevski, J. Pal, R. Patra, S. Surana, S. Surana, and K. Fall, “The Case of Technology for Developing Regions,” *Computer*, vol. 38, no. 6, pp. 25–38, June 2005.
- [12] B. Du, M. Demmer, and E. Brewer, “Analysis of WWW Traffic in Cambodia and Ghana,” *Proc. Int'l World Wide Web (WWW '06) Conf.*, 2006.
- [13] P. Juang, H. Oki, Y. Wang, M. Martonosi, L.S. Peh, and D. Rubenstein, “Energy-Efficient Computing for Wildlife Tracking: Design Tradeoffs and Early Experiences with Zebnet,” *Proc. ACM Int'l Conf. Architectural Support for Programming Languages and Operating Systems (ASPLOS '02)*, 2002.
- [14] E.P. Jones and P.A. Ward, “Routing Strategies for Delay-Tolerant Networks,” <http://cng.uwaterloo.ca/pasward/Publications/>, 2007.
- [15] T. Spyropoulos, K. Psounis, and C.S. Raghavendra, “Efficient Routing in Intermittently Connected Mobile Networks: The Single-Copy Case,” *IEEE/ACM Trans. Networking*, vol. 16, no. 1, pp. 63–76, Feb. 2008.
- [16] T. Spyropoulos, K. Psounis, and C.S. Raghavendra, “Efficient Routing in Intermittently Connected Mobile Networks: The Multiple-Copy Case,” *IEEE/ACM Trans. Networking*, vol. 16, no. 1, pp. 77–90, Feb. 2008.
- [17] T. Small and Z. Haas, “Resource and Performance Tradeoffs in Delay-Tolerant Wireless Networks,” *Proc. ACM SIGCOMM Workshop Delay Tolerant Networking (WDTN '05)*, 2005.
- [18] G. Neglia and X. Zhang, “Optimal Delay-Power Tradeoff in Sparse Delay Tolerant Networks: A Preliminary Study,” *Proc. ACM SIGCOMM Workshop Challenged Networks (CHANTS '06)*, 2006.
- [19] A. Vahdat and D. Becker, “Epidemic Routing for Partially Connected Ad Hoc Networks,” Technical Report CS-200006, Duke Univ., 2000.
- [20] A. Lindgren, A. Doria, and O. Schelen, “Probabilistic Routing in Intermittently Connected Networks,” *SIGMOBILE Mobile Computing and Comm. Rev.*, vol. 7, no. 3, 2003.
- [21] J. Burgess, G. Bissias, M. Corner, and B. Levine, “Surviving Attacks on Disruption-Tolerant Networks without Authentication,” *Proc. ACM MobiHoc*, 2007.
- [22] T. Spyropoulos, K. Psounis, and C.S. Raghavendra, “Spray and Wait: Efficient Routing in Intermittently Connected Mobile Networks,” *Proc. ACM SIGCOMM Workshop Delay Tolerant Networking (WDTN '05)*, 2005.

- [23] H. Dubois-Ferriere, M. Grossglauser, and M. Vetterli, "Age Matters: Efficient Route Discovery in Mobile Ad Hoc Networks Using Encounter Ages," *Proc. ACM MobiHoc*, 2003.
- [24] S. Zhong, J. Chen, and Y.R. Yang, "Sprite: A Simple, Cheat-Proof, Credit-Based System for Mobile Ad-Hoc Networks," *Proc. IEEE INFOCOM*, 2003.
- [25] T. Spyropoulos, T. Turletti, and K. Obraczka, "Utility-Based Message Replication for Intermittently Connected Heterogeneous Wireless Networks," *Proc. IEEE World of Wireless, Mobile and Multimedia Networks (WoWMoM) Workshop Autonomic and Opportunistic Comm. (AOC '07)*, 2007.
- [26] D.J. Daley and J. Gani, *Epidemic Modelling: An Introduction*. Cambridge Univ. Press, 2001.
- [27] X. Zhang, G. Neglia, J. Kurose, and D. Towsley, "Performance Modeling of Epidemic Routing," *Proc. Int'l Federation for Information Processing (IFIP) Networking Conf.*, 2006.
- [28] V. Lenders, G. Karlsson, and M. May, "Wireless Ad Hoc Podcasting," *Proc. IEEE Conf. Sensor and Ad Hoc Comm. and Networks (SECON '07)*, 2007.
- [29] W. Zhao, M. Ammar, and E. Zegura, "A Message Ferrying Approach for Data Delivery in Sparse Mobile Ad Hoc Networks," *Proc. ACM MobiHoc*, 2004.
- [30] M.S.A. Hall and S. Hippler, "Multicommodity Flows over Time: Efficient Algorithms and Complexity," *Proc. Int'l Colloquium on Automata, Languages and Programming (ICALP '03)*, 2003.
- [31] B. Hoppe and E. Tardos, "The Quickest Transshipment Problem," *Proc. Symp. Discrete Algorithms (SODA '95)*, 1995.
- [32] E.P.C. Jones, L. Li, and P.A.S. Ward, "Practical Routing in Delay-Tolerant Networks," *Proc. ACM SIGCOMM Workshop Delay Tolerant Networking (WDTN '05)*, 2005.
- [33] N. Sarafijanovic-Djukic, M. Piorkowski, and M. Grossglauser, "Island Hopping: Efficient Mobility-Assisted Forwarding in Partitioned Networks," *Proc. IEEE Conf. Sensor and Ad Hoc Comm. and Networks (SECON '06)*, 2006.
- [34] J. Ott, D. Kutscher, and C. Dwertmann, "Integrating DTN and MANET Routing," *Proc. SIGCOMM Workshop Challenged Networks (CHANTS '06)*, 2006.
- [35] A. Chaintreau, P. Hui, J. Crowcroft, C. Diot, R. Gass, and J. Scott, "Impact of Human Mobility on the Design of Opportunistic Forwarding Algorithms," *Proc. IEEE INFOCOM*, 2006.
- [36] J. Burgess, B. Gallagher, D. Jensen, and B.N. Levine, "MaxProp: Routing for Vehicle-Based Disruption-Tolerant Networks," *Proc. IEEE INFOCOM*, 2006.
- [37] J. Zhao and G. Cao, "Vadd: Vehicle-Assisted Data Delivery in Vehicular Ad Hoc Networks," *Proc. IEEE INFOCOM*, 2006.
- [38] S. Jain, M. Demmer, R. Patra, and K. Fall, "Using Redundancy to Cope with Failures in a Delay Tolerant Network," *Proc. ACM SIGCOMM*, 2005.
- [39] A. Balasubramanian, B. Levine, and A. Venkataramani, "DTN Routing as a Resource Allocation Problem," *Proc. ACM SIGCOMM*, 2007.
- [40] Z.J. Haas and T. Small, "A New Networking Model for Biological Applications of Ad Hoc Sensor Networks," *IEEE/ACM Trans. Networking*, vol. 14, no. 1, pp. 27-40, Feb. 2006.
- [41] R. Groenevelt, G. Koole, and P. Nain, "Message Delay in Manet (Extended Abstract)," *Proc. ACM Sigmetrics*, 2005.
- [42] A. Jindal and K. Psounis, "Performance Analysis of Epidemic Routing under Contention," *Proc. ACM Int'l Wireless Comm. and Mobile Computing Conf. (IWCMC '06)*, 2006.
- [43] F. Ball and O.D. Lyne, "Stochastic Multitype SIR Epidemics among a Population Partitioned into Households," *Advances in Applied Probability*, vol. 33, no. 1, pp. 99-123, 2001.
- [44] R.K. Watson, "On an Epidemic in a Stratified Population," *J. Applied Probability*, vol. 9, no. 3, pp. 659-666, 1972.
- [45] S. Rushton and A.J. Mautner, "The Deterministic Model of a Simple Epidemic for More than one Community," *Biometrika*, vol. 42, nos. 1/2, pp. 126-132, 1955.
- [46] C. Moore and M.E.J. Newman, "Epidemics and Percolation in Small-World Networks," *Physical Rev. E*, vol. 61, p. 5678, 2000.
- [47] R. Durrett, *Probability: Theory and Examples*, second ed. Duxbury Press, 1995.
- [48] D. Aldous and J. Fill, "Reversible Markov Chains and Random Walks on Graphs," monograph in preparation, <http://stat-www.berkeley.edu/users/aldous/RWG/book.html>, 2009.
- [49] A. Jindal and K. Psounis, "Optimizing Multi-Copy Routing Schemes for Resource Constrained Intermittently Connected Mobile Networks," *Proc. IEEE Asilomar Conf.*, 2006.
- [50] M. Musolesi and C. Mascolo, "A Community Based Mobility Model for Ad Hoc Network Research," *Proc. ACM REALMAN Workshop*, 2006.
- [51] J. Leguay, T. Friedman, and V. Conan, "DTN Routing in a Mobility Pattern Space," *Proc. ACM SIGCOMM Workshop Delay Tolerant Networking (WDTN '05)*, 2005.
- [52] P. Hui and J. Crowcroft, "How Small Labels Create Big Improvements," *Proc. IEEE Int'l Workshop Intermittently Connected Mobile Ad hoc Networks (ICMAN '07)*, 2007.
- [53] The UMASS DieselNet Project, <http://prisms.cs.umass.edu/dome/index.php?page=umassdieselnet>, 2009.
- [54] CRAWDAD: A Community Resource for Archiving Wireless Data at Dartmouth, <http://crawdad.cs.dartmouth.edu/>, 2009.
- [55] DTN Java Simulator, <http://people.ee.ethz.ch/~spyropot/dtnsim.html>, 2009.
- [56] T. Spyropoulos, K. Psounis, and C.S. Raghavendra, "Performance Analysis of Mobility-Assisted Routing," *Proc. ACM MobiHoc*, 2006.
- [57] W. Hsu, T. Spyropoulos, K. Psounis, and A. Helmy, "Modeling Time-Variant User Mobility in Wireless Mobile Networks," *Proc. IEEE INFOCOM*, 2007.
- [58] W. Hsu and A. Helmy, "IMPACT: Investigation of Mobile-User Patterns Across University Campuses Using WLAN Trace Analysis," technical report, Univ. of Southern California, 2005.
- [59] G.J. Alexander and J.C. Francis, *Portfolio Analysis*. Prentice Hall, 1986.
- [60] A. Seth and S. Keshav, "Practical Security for Disconnected Nodes," *Proc. Int'l Conf. Network Protocols (ICNP '05) Workshop Secure Network Protocols (NPsec)*, 2005.
- [61] M. Chuah, P. Yang, and J. Han, "A Ferry-Based Intrusion Detection Scheme for Sparsely Connected Adhoc Networks," *Proc. First Workshop Security for Emerging Ubiquitous Computing*, 2007.
- [62] A. El Fawal, J.-Y. Le Boudec, and K. Salamatian, "Vulnerabilities of Epidemic Forwarding," *Proc. IEEE World of Wireless, Mobile and Multimedia Networks (WoWMoM '07) Workshop Autonomic and Opportunistic Comm. (AOC)*, 2007.
- [63] S. Capkun, J.-P. Hubaux, and L. Buttyán, "Mobility Helps Security in Ad Hoc Networks," *Proc. ACM MobiHoc*, 2003.
- [64] D. Dagon, T. Martin, and T. Starner, "Mobile Phones as Computing Devices: The Viruses Are Coming!" *IEEE Pervasive Computing*, vol. 3, no. 4, pp. 11-15, Oct.-Dec. 2004.



He is a member of the IEEE.



Group at INRIA, Sophia Antipolis. He is a member of the IEEE.



Thrasyvoulos Spyropoulos received the diploma in electrical and computer engineering from the National Technical University of Athens, Greece, in February 2000, and the PhD degree in electrical engineering from the University of Southern California in May 2006. From 2006 to 2007, he was a postdoctoral researcher at INRIA, Sophia-Antipolis. He is currently a senior researcher with the Swiss Federal Institute of Technology (ETH), Zürich.

Thierry Turletti received the BE, MS (1990), and PhD (1995) degrees in computer science from the University of Nice—Sophia Antipolis, France. His PhD work was performed with the RODEO project at INRIA, Sophia-Antipolis. From 1995 to 1996, he was a postdoctoral fellow in the Telemedia, Networks, and Systems Group at the Massachusetts Institute of Technology's Laboratory for Computer Science. He is currently a full-time researcher in the Planete

Katia Obraczka received the BS and MS degrees in electrical and computer engineering from the Federal University of Rio de Janeiro. She received the MS and PhD degrees in computer science from the University of Southern California in 1990 and 1994, respectively. She is currently an associate professor at the University of California, Santa Cruz's Jack Baskin School of Engineering. She is a member of the IEEE.